

Capstone Project 9. Comprehensive Data Pipeline with Airflow, Kafka, PySpark and Visualization on for Credit Card Approval Analysis.

Name: Varsha .S

Roll On: 23EE056

Objective:

The objective of this capstone project is to build a robust, scalable, and fault-tolerant end-to-end data pipeline for an e-commerce platform. The system will ingest real-time order and customer data using Kafka, transform and process this data using PySpark, and store the results in a queryable format. The transformed data will be visualized using Plotly or Power BI to generate actionable insights. The entire workflow will be orchestrated and automated using Apache Airflow, enabling scheduled execution, monitoring, and dependency management.

This project demonstrates the ability to design modern data pipelines combining streaming ingestion, batch processing, distributed transformation, orchestration, and business intelligence reporting in a production-ready architecture.

Objectives:

1. Data Extraction: Ingest data from multiple sources using HDFS.
2. Data Transformation: Use PySpark in for data cleaning, transformation, and enrichment.
3. Materialized Views: Create SQL-based materialized views for efficient querying and analysis.
4. Visualization: Develop a visualization layer using Power BI or Plotly to present the Visualization.

Architecture:

1. Input Data Sources (Multiple):
 - Use Data sources such as HDFS, and relational databases (MySQL).
 - Read dataset from multiple sources using Pyspark or Stream Data using Kafka Producer.
2. Materialized Views (SQL):
 - Transform data using PySpark.

- Load the transformed data into SQL Database.
- Create materialized views for optimized querying and reporting.

3. Visualization Layer (Power BI/Dashboard):

- Create interactive visualization of analysis to visualize key metrics and insights from dataset.

Dataset:

https://www.kaggle.com/datasets/rikdifos/credit-card-approval-prediction?select=credit_record.csv

Dataset 1

- ID: Unique identifier for each client.
- CODE_GENDER: Gender of the client.
- FLAG_OWN_CAR: Whether the client owns a car.
- FLAG_OWN_REALTY: Whether the client owns property.
- CNT_CHILDREN: The number of children the client has.
- AMT_INCOME_TOTAL: The total annual income of the client.
- NAME_INCOME_TYPE: The category of the client's income source.
- NAME_EDUCATION_TYPE: The highest education level the client has achieved.
- NAME_FAMILY_STATUS: The marital status of the client.
- NAME_HOUSING_TYPE: The client's living situation.
- DAYS_BIRTH: The client's age in days, counted backwards from the current day.
- DAYS_EMPLOYED: How long the client has been employed, counted backwards from the current day. Positive numbers indicate unemployment.
- FLAG_MOBIL: Whether the client owns a mobile phone.
- FLAG_WORK_PHONE: Whether the client has a work phone.
- FLAG_PHONE: Whether the client has a phone.
- FLAG_EMAIL: Whether the client has an email address.
- OCCUPATION_TYPE: The client's occupation.
- CNT_FAM_MEMBERS: The size of the client's family.

The credit_record.csv dataset tracks the credit history of clients, with each record reflecting a monthly snapshot of an individual's credit file.

Dataset 2

- ID: Unique identifier for each client, matching the ID in the application_record.csv.
- MONTHS_BALANCE: The month of the record relative to the current month (0 is current, -1 is previous month, etc.).
- STATUS: The status of the client's credit for that month (e.g., no overdue, days past due, paid off).

Steps to Implement:

1. Data Ingestion Layer (Kafka)

Goal:

Ingest real-time data such as customer orders, product views, cart additions, etc., into the pipeline.

Steps:

- Use a Python Kafka producer to send these simulated events to Kafka topics
- Set up Kafka topics and configure Kafka brokers and Zookeeper locally or on a cluster.
- Use Kafka Consumer within Pyspark Structured Streaming to read incoming data in near real-time.

Perform transformations such as:

- Join application + credit records on ID.
- Aggregate credit STATUS by MONTHS_BALANCE.
- Show trends in loan repayment status (e.g., counts of STATUS = 0 = no DPD, 1 = 1–30 DPD, etc.).

Technologies:

Kafka, Python Producer, JSON/CSV format

Outputs:

```
hadoop@hadoop-VirtualBox: ~  
File Edit View Search Terminal Help  
hadoop@hadoop-VirtualBox:~$ start-dfs.sh  
Starting namenodes on [localhost]  
localhost: starting namenode, logging to /usr/local/hadoop/logs/hadoop-hadoop-namenode-hadoop-VirtualBox.out  
localhost: starting datanode, logging to /usr/local/hadoop/logs/hadoop-hadoop-datanode-hadoop-VirtualBox.out  
Starting secondary namenodes [0.0.0.0]  
0.0.0.0: starting secondarynamenode, logging to /usr/local/hadoop/logs/hadoop-hadoop-secondarynamenode-hadoop-VirtualBox.out  
hadoop@hadoop-VirtualBox:~$ start-yarn.sh  
starting yarn daemons  
starting resourcemanager, logging to /usr/local/hadoop/logs/yarn-hadoop-resourcemanager-hadoop-VirtualBox.out  
localhost: starting nodemanager, logging to /usr/local/hadoop/logs/yarn-hadoop-nodemanager-hadoop-VirtualBox.out  
hadoop@hadoop-VirtualBox:~$ jps  
3094 Jps  
2123 NameNode  
2685 ResourceManager  
2270 DataNode  
2511 SecondaryNameNode  
2831 NodeManager
```

```
hadoop@hadoop-VirtualBox:~$ hdfs dfs -mkdir /user/hadoop/credit_project
```

```
hadoop@hadoop-VirtualBox:~$ hdfs dfs -put /home/hadoop/application_record.csv /user/hadoop/credit_project  
hadoop@hadoop-VirtualBox:~$ hdfs dfs -put /home/hadoop/credit_record.csv /user/hadoop/credit_project/  
hadoop@hadoop-VirtualBox:~$ hdfs dfs -ls /user/hadoop/credit_project  
Found 2 items  
-rw-r--r-- 1 hadoop supergroup 53028297 2025-09-14 11:16 /user/hadoop/credit_project/application_record.csv  
-rw-r--r-- 1 hadoop supergroup 15367102 2025-09-14 11:17 /user/hadoop/credit_project/credit_record.csv
```

START ZOOKEEPER

```
hadoop@hadoop-VirtualBox:~/Downloads$ cd ~/Downloads/kafka_2.13-3.7.0
hadoop@hadoop-VirtualBox:~/Downloads/kafka_2.13-3.7.0$ bin/zookeeper-server-start.sh config/zookeeper.properties
[2025-09-14 11:19:58,028] INFO Reading configuration from: config/zookeeper.properties (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2025-09-14 11:19:58,032] WARN config/zookeeper.properties is relative. Prepend ./ to indicate that you're sure! (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2025-09-14 11:19:58,105] INFO clientPortAddress is 0.0.0.0:2181 (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2025-09-14 11:19:58,106] INFO secureClientPort is not set (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2025-09-14 11:19:58,106] INFO observerMasterPort is not set (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2025-09-14 11:19:58,106] INFO metricsProvider.className is org.apache.zookeeper.metrics.impl.DefaultMetricsProvider (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2025-09-14 11:19:58,117] INFO autopurge.snapRetainCount set to 3 (org.apache.zookeeper.server.DataDirCleanupManager)
[2025-09-14 11:19:58,118] INFO autopurge.purgeInterval set to 0 (org.apache.zookeeper.server.DataDirCleanupManager)
[2025-09-14 11:19:58,118] INFO Purge task is not scheduled. (org.apache.zookeeper.server.DataDirCleanupManager)
[2025-09-14 11:19:58,118] WARN Either no config or no quorum defined in config, running in standalone mode (org.apache.zookeeper.server.quorum.QuorumPeerMain)
[2025-09-14 11:19:58,123] INFO Log4j 1.2 jmx support not found; jmx disabled. (org.apache.zookeeper.jmx.ManagedUtil)
[2025-09-14 11:19:58,127] INFO Reading configuration from: config/zookeeper.properties (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2025-09-14 11:19:58,127] WARN config/zookeeper.properties is relative. Prepend ./ to indicate that you're sure! (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2025-09-14 11:19:58,128] INFO clientPortAddress is 0.0.0.0:2181 (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2025-09-14 11:19:58,129] INFO secureClientPort is not set (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2025-09-14 11:19:58,129] INFO observerMasterPort is not set (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2025-09-14 11:19:58,129] INFO metricsProvider.className is org.apache.zookeeper.metrics.impl.DefaultMetricsProvider (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2025-09-14 11:19:58,129] INFO Starting server (org.apache.zookeeper.server.ZooKeeperServerMain)
[2025-09-14 11:19:58,252] INFO ServerMetrics initialized with provider org.apache.zookeeper.metrics.impl.DefaultMetricsProvider@327471b5 (org.apache.zookeeper.server.ServerMetrics)
[2025-09-14 11:19:58,267] INFO ACL digest algorithm is: SHA1 (org.apache.zookeeper.server.auth.DigestAuthenticationProvider)
[2025-09-14 11:19:58,267] INFO zookeeper.DigestAuthenticationProvider.enabled = true (org.apache.zookeeper.server.auth.DigestAuthenticationProvider)
```

START KAFKA BROKER


```
from kafka import KafkaProducer
import pandas as pd
import json, time

# Load CSVs
app_df = pd.read_csv("/home/hadoop/Downloads/application_record.csv")
credit_df = pd.read_csv("/home/hadoop/Downloads/credit_record.csv")

# Merge datasets
merged = pd.merge(credit_df, app_df, on="ID", how="inner")

# Kafka producer
producer = KafkaProducer(
    bootstrap_servers=['localhost:9092'],
    value_serializer=lambda v: json.dumps(v).encode('utf-8')
)

# Stream row by row
for _, row in merged.iterrows():
    event = row.to_dict()
    producer.send("credit_events", value=event)
    print("Sent:", event)
    time.sleep(1)
```



```

d.csv
hadoop@hadoop-VirtualBox:~$ nano producer.py
hadoop@hadoop-VirtualBox:~$ python3 producer.py
Sent: {'ID': 5008804, 'MONTHS_BALANCE': 0, 'STATUS': 'C', 'CODE_GENDER': 'M', 'FLAG_OWN_CAR': 'Y', 'FLAG_OWN_REALTY': 'Y', 'CNT_CHILDREN': 0, 'AMT_INCOME_TOTAL': 427500.0, 'NAME_INCOME_TYPE': 'Working', 'NAME_EDUCATION_TYPE': 'Higher education', 'NAME_FAMILY_STATUS': 'Civil marriage', 'NAME_HOUSING_TYPE': 'Rented apartment', 'DAYS_BIRTH': -12005, 'DAYS_EMPLOYED': -4542, 'FLAG_MOBIL': 1, 'FLAG_WORK_PHONE': 1, 'FLAG_PHONE': 0, 'FLAG_EMAIL': 0, 'OCCUPATION_TYPE': nan, 'CNT_FAM_MEMBERS': 2}
Sent: {'ID': 5008804, 'MONTHS_BALANCE': -1, 'STATUS': 'C', 'CODE_GENDER': 'M', 'FLAG_OWN_CAR': 'Y', 'FLAG_OWN_REALTY': 'Y', 'CNT_CHILDREN': 0, 'AMT_INCOME_TOTAL': 427500.0, 'NAME_INCOME_TYPE': 'Working', 'NAME_EDUCATION_TYPE': 'Higher education', 'NAME_FAMILY_STATUS': 'Civil marriage', 'NAME_HOUSING_TYPE': 'Rented apartment', 'DAYS_BIRTH': -12005, 'DAYS_EMPLOYED': -4542, 'FLAG_MOBIL': 1, 'FLAG_WORK_PHONE': 1, 'FLAG_PHONE': 0, 'FLAG_EMAIL': 0, 'OCCUPATION_TYPE': nan, 'CNT_FAM_MEMBERS': 2}
Sent: {'ID': 5008804, 'MONTHS_BALANCE': -2, 'STATUS': 'C', 'CODE_GENDER': 'M', 'FLAG_OWN_CAR': 'Y', 'FLAG_OWN_REALTY': 'Y', 'CNT_CHILDREN': 0, 'AMT_INCOME_TOTAL': 427500.0, 'NAME_INCOME_TYPE': 'Working', 'NAME_EDUCATION_TYPE': 'Higher education', 'NAME_FAMILY_STATUS': 'Civil marriage', 'NAME_HOUSING_TYPE': 'Rented apartment', 'DAYS_BIRTH': -12005, 'DAYS_EMPLOYED': -4542, 'FLAG_MOBIL': 1, 'FLAG_WORK_PHONE': 1, 'FLAG_PHONE': 0, 'FLAG_EMAIL': 0, 'OCCUPATION_TYPE': nan, 'CNT_FAM_MEMBERS': 2}
Sent: {'ID': 5008804, 'MONTHS_BALANCE': -3, 'STATUS': 'C', 'CODE_GENDER': 'M', 'FLAG_OWN_CAR': 'Y', 'FLAG_OWN_REALTY': 'Y', 'CNT_CHILDREN': 0, 'AMT_INCOME_TOTAL': 427500.0, 'NAME_INCOME_TYPE': 'Working', 'NAME_EDUCATION_TYPE': 'Higher education', 'NAME_FAMILY_STATUS': 'Civil marriage', 'NAME_HOUSING_TYPE': 'Rented apartment', 'DAYS_BIRTH': -12005, 'DAYS_EMPLOYED': -4542, 'FLAG_MOBIL': 1, 'FLAG_WORK_PHONE': 1, 'FLAG_PHONE': 0, 'FLAG_EMAIL': 0, 'OCCUPATION_TYPE': nan, 'CNT_FAM_MEMBERS': 2}
Sent: {'ID': 5008804, 'MONTHS_BALANCE': -4, 'STATUS': 'C', 'CODE_GENDER': 'M', 'FLAG_OWN_CAR': 'Y', 'FLAG_OWN_REALTY': 'Y', 'CNT_CHILDREN': 0, 'AMT_INCOME_TOTAL': 427500.0, 'NAME_INCOME_TYPE': 'Working', 'NAME_EDUCATION_TYPE': 'Higher education', 'NAME_FAMILY_STATUS': 'Civil marriage', 'NAME_HOUSING_TYPE': 'Rented apartment', 'DAYS_BIRTH': -12005, 'DAYS_EMPLOYED': -4542, 'FLAG_MOBIL': 1, 'FLAG_WORK_PHONE': 1, 'FLAG_PHONE': 0, 'FLAG_EMAIL': 0, 'OCCUPATION_TYPE': nan, 'CNT_FAM_MEMBERS': 2}
Sent: {'ID': 5008804, 'MONTHS_BALANCE': -5, 'STATUS': 'C', 'CODE_GENDER': 'M', 'FLAG_OWN_CAR': 'Y', 'FLAG_OWN_REALTY': 'Y', 'CNT_CHILDREN': 0, 'AMT_INCOME_TOTAL': 427500.0, 'NAME_INCOME_TYPE': 'Working', 'NAME_EDUCATION_TYPE': 'Higher education', 'NAME_FAMILY_STATUS': 'Civil marriage', 'NAME_HOUSING_TYPE': 'Rented apartment', 'DAYS_BIRTH': -12005, 'DAYS_EMPLOYED': -4542, 'FLAG_MOBIL': 1, 'FLAG_WORK_PHONE': 1, 'FLAG_PHONE': 0, 'FLAG_EMAIL': 0, 'OCCUPATION_TYPE': nan, 'CNT_FAM_MEMBERS': 2}
Sent: {'ID': 5008804, 'MONTHS_BALANCE': -6, 'STATUS': 'C', 'CODE_GENDER': 'M', 'FLAG_OWN_CAR': 'Y', 'FLAG_OWN_REALTY': 'Y', 'CNT_CHILDREN': 0, 'AMT_INCOME_TOTAL': 427500.0, 'NAME_INCOME_TYPE': 'Working', 'NAME_EDUCATION_TYPE': 'Higher education', 'NAME_FAMILY_STATUS': 'Civil marriage', 'NAME_HOUSING_TYPE': 'Rented apartment', 'DAYS_BIRTH': -12005, 'DAYS_EMPLOYED': -4542, 'FLAG_MOBIL': 1, 'FLAG_WORK_PHONE': 1, 'FLAG_PHONE': 0, 'FLAG_EMAIL': 0, 'OCCUPATION_TYPE': nan, 'CNT_FAM_MEMBERS': 2}

```

Jupyter Notebook In VM

```

from pyspark.sql import SparkSession

spark = SparkSession.builder \
    .appName("CreditConsumer") \
    .getOrCreate()

```



```

In [2]: from pyspark.sql.functions import from_json, col
        from pyspark.sql.types import StructType, StringType, IntegerType, DoubleType

        # Schema for merged dataset
        schema = StructType() \
            .add("ID", StringType()) \
            .add("MONTHS_BALANCE", IntegerType()) \
            .add("STATUS", StringType()) \
            .add("CODE_GENDER", StringType()) \
            .add("FLAG_OWN_CAR", StringType()) \
            .add("FLAG_OWN_REALTY", StringType()) \
            .add("CNT_CHILDREN", IntegerType()) \
            .add("AMT_INCOME_TOTAL", DoubleType()) \
            .add("NAME_INCOME_TYPE", StringType()) \
            .add("NAME_EDUCATION_TYPE", StringType()) \
            .add("NAME_FAMILY_STATUS", StringType()) \
            .add("NAME_HOUSING_TYPE", StringType()) \
            .add("DAYS_BIRTH", IntegerType()) \
            .add("DAYS_EMPLOYED", IntegerType()) \
            .add("OCCUPATION_TYPE", StringType()) \
            .add("CNT_FAM_MEMBERS", DoubleType())

        # Read from Kafka
        df = spark.readStream.format("kafka") \
            .option("kafka.bootstrap.servers", "localhost:9092") \
            .option("subscribe", "credit_events") \
            .load()

        # Parse JSON
        value_df = df.selectExpr("CAST(value AS STRING)") \
            .select(from_json(col("value"), schema).alias("data")) \
            .select("data.*")

        value_df.printSchema()

root
|-- ID: string (nullable = true)
|-- MONTHS_BALANCE: integer (nullable = true)
|-- STATUS: string (nullable = true)
|-- CODE_GENDER: string (nullable = true)
|-- FLAG_OWN_CAR: string (nullable = true)

```

BATCH STREAMING PROCESS

```

In [3]: # Count repayment status categories
        status_count = value_df.groupBy("STATUS").count()

        # Start streaming query to console
        query = status_count.writeStream \
            .outputMode("complete") \
            .format("console") \
            .start()

```

```
hadoop@hadoo... x hadoop@hadoo... x hadoop@hadoo... x hadoop@hadoo... x hadoop@hadoo... x
Batch: 1
-----
+-----+-----+
|STATUS|count|
+-----+-----+
|      0|   13|
|      C|   12|
|      X|    1|
|      1|    2|
+-----+-----+

[Stage 5:=====> (29 + 2) / 200]25/09/14 11:34:13
WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1245)
    at java.lang.Thread.join(Thread.java:1319)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.java:609)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:370)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:546)
[Stage 5:=====> (151 + 2) / 200]25/09/14 11:34:27
WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1245)
    at java.lang.Thread.join(Thread.java:1319)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.java:609)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeInternal(DFSOutputStream.java:577)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:573)
-----
Batch: 2
-----
+-----+-----+
|STATUS|count|
+-----+-----+
|      0|   13|
|      C|   29|
|      X|    1|
|      1|    2|
+-----+-----+

[Stage 7:=====> (171 + 2) / 200]
```

2. Data Transformation Layer (PySpark)

Goal:

Cleanse, enrich, and join raw data to create meaningful datasets for downstream analytics.

Steps:

- Read streaming data from Kafka using Pyspark Structured Streaming.
- Parse incoming JSON or CSV payloads and convert to Dataframes.
- Perform transformations:
 - Data type casting
 - Handling missing values

- Joining with static datasets (if multiple datasets are present)
- Data Aggregations if required
- Time-windowed operations (e.g. Rides Booking)
- Write transformed data into:
 - Delta Lake / HDFS / Parquet format for historical records
 - Temporary views for reporting

Technologies:

PySpark, HDFS, Structured Streaming

Outputs:

```
In [4]: from pyspark.sql import SparkSession
from pyspark.sql.functions import from_json, col
from pyspark.sql.types import StructType, StructField, StringType, IntegerType

# Spark session with Kafka support
spark = SparkSession.builder \
    .appName("CreditDataTransformation") \
    .config("spark.sql.streaming.checkpointLocation", "/tmp/checkpoints") \
    .getOrCreate()

# Read stream from Kafka
df_raw = spark.readStream \
    .format("kafka") \
    .option("kafka.bootstrap.servers", "localhost:9092") \
    .option("subscribe", "credit_events") \
    .load()

# Extract value from Kafka (payload)
df_raw = df_raw.selectExpr("CAST(value AS STRING)")
```

```
In [5]: schema_app = StructType([
    StructField("ID", StringType()),
    StructField("CODE_GENDER", StringType()),
    StructField("FLAG_OWN_CAR", StringType()),
    StructField("FLAG_OWN_REALTY", StringType()),
    StructField("CNT_CHILDREN", IntegerType()),
    StructField("AMT_INCOME_TOTAL", DoubleType()),
    StructField("NAME_INCOME_TYPE", StringType()),
    StructField("NAME_EDUCATION_TYPE", StringType()),
    StructField("NAME_FAMILY_STATUS", StringType()),
    StructField("NAME_HOUSING_TYPE", StringType()),
    StructField("DAYS_BIRTH", IntegerType()),
    StructField("DAYS_EMPLOYED", IntegerType()),
    StructField("FLAG_MOBIL", IntegerType()),
    StructField("FLAG_WORK_PHONE", IntegerType()),
    StructField("FLAG_PHONE", IntegerType()),
    StructField("FLAG_EMAIL", IntegerType()),
```

AFTER PARSING, CLEANING AND STORING IT IN PARQUET IN HDFS=>

```
In [24]: df_joined.writeStream \
        .format("parquet") \
        .option("path", "hdfs://localhost:9000/user/hadoop/credit_project/out
        .option("checkpointLocation", "/tmp/checkpoints/credit_joined") \
        .outputMode("append") \
        .start()
```

```
Out[24]: <pyspark.sql.streaming.StreamingQuery at 0x7f30a85debe0>
```

```
In [*]: df_joined.createOrReplaceTempView("credit_analysis")

result = spark.sql("""
    SELECT STATUS, AVG(AMT_INCOME_TOTAL) as avg_income
    FROM credit_analysis
    GROUP BY STATUS
""")
```

```
In [26]: df_output = spark.read.parquet("hdfs://localhost:9000/user/hadoop/credit_
df_output.show(10, truncate=False)
```

```
+---+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|ID |MONTHS_BALANCE|STATUS|CODE |GENDER|FLAG_OWN_CAR|FLAG_OWN_REALTY|CNT_
CHILDREN|AMT_INCOME_TOTAL|NAME_INCOME_TYPE|NAME_EDUCATION_TYPE|NAME_FAMI
LY_STATUS|NAME_HOUSING_TYPE|DAYS_BIRTH|DAYS_EMPLOYED|FLAG_MOBIL|FLAG_WOR
K_PHONE|FLAG_PHONE|FLAG_EMAIL|OCCUPATION_TYPE|CNT_FAM_MEMBERS|
+---+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+---+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
```

3. Creation of Materialized Views

Now that our data is transformed and stored, we will create materialized views for efficient querying.

- **Define Views:** Identify key metrics and insights that the business needs. Define SQL queries to create materialized views based on these requirements.
- **Optimization:** Optimize these views for performance, ensuring they are quick to query and provide accurate results.

- Use insightful queries to find good insights combining all data together.
- Calculate bad customer rate
- Survival Analysis
- Cumulative % of Bad Customers Analysis
- Observe Window Analysis

Tools – Pyspark, Spark SQL

Outputs:

```
app_df = spark.read.csv("/user/hadoop/credit_project/application_record.csv")
# Load credit_record
credit_df = spark.read.csv("/user/hadoop/credit_project/credit_record.csv")
# Quick check
app_df.show(5)
credit_df.show(5)
```

```
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
|      ID|CODE_GENDER|FLAG_OWN_CAR|FLAG_OWN_REALTY|CNT_CHILDREN|AMT_INCOM
E_TOTAL|      NAME_INCOME_TYPE|      NAME_EDUCATION_TYPE|      NAME_FAMILY_STATUS|N
AME_HOUSING_TYPE|DAYS_BIRTH|DAYS_EMPLOYED|FLAG_MOBIL|FLAG_WORK_PHONE|FLA
G_PHONE|FLAG_EMAIL|OCCUPATION_TYPE|CNT_FAM_MEMBERS|
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
|5008804|      M|      Y|      Y|      0|      4
27500.0|      Working|      Higher education|      Civil marriage|
Rented apartment|      -12005|      -4542|      1|      1|
0|      0|      null|      2|
|5008805|      M|      Y|      Y|      0|      4
27500.0|      Working|      Higher education|      Civil marriage|
Rented apartment|      -12005|      -4542|      1|      1|
0|      0|      null|      2|
|5008806|      M|      Y|      Y|      0|      1
12500.0|      Working|Secondary / secon...|      Married|H
ouse / apartment|      -21474|      -1134|      1|      0|
0|      0|Security staff|      2|
|5008808|      F|      N|      Y|      0|      2
70000.0|Commercial associate|Secondary / secon...|Single / not married|H
ouse / apartment|      -19110|      -3051|      1|      0|
1|      1|Sales staff|      1|
|5008809|      F|      N|      Y|      0|      2
70000.0|Commercial associate|Secondary / secon...|Single / not married|H
ouse / apartment|      -19110|      -3051|      1|      0|
```

```

+|          +|      sales status|          +|
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+

```

only showing top 5 rows

```

+-----+-----+-----+
|      ID|MONTHS_BALANCE|STATUS|
+-----+-----+-----+
|5001711|          0|    X|
|5001711|         -1|    0|
|5001711|         -2|    0|
|5001711|         -3|    0|
|5001712|          0|    C|
+-----+-----+-----+

```

only showing top 5 rows

```

In [39]: spark.sql("""
SELECT
    COUNT(*) AS total_customers,
    SUM(BAD_FLAG) AS bad_customers,
    ROUND(SUM(BAD_FLAG)/COUNT(*)*100,2) AS bad_customer_rate
FROM credit_analysis_view
""").show()

```

```

+-----+-----+-----+
|total_customers|bad_customers|bad_customer_rate|
+-----+-----+-----+
|          777715|          11575|             1.49|
+-----+-----+-----+

```

Bad Customer

```
In [40]: from pyspark.sql.window import Window
from pyspark.sql.functions import col, sum as _sum, desc

windowSpec = Window.orderBy(desc("BAD_FLAG")).rowsBetween(Window.unboundedPreceding, Window.unboundedFollowing)

cum_df = joined_df.withColumn("cum_bad", _sum("BAD_FLAG").over(windowSpec))
                  .withColumn("cum_bad_percent", col("cum_bad")/_sum("BAD_FLAG").over(windowSpec))

cum_df.select("ID", "BAD_FLAG", "cum_bad", "cum_bad_percent").show(10)
```

ID	BAD_FLAG	cum_bad	cum_bad_percent
5008804	1	1	0.008639308855291577
5008805	1	2	0.017278617710583154
5008825	1	3	0.02591792656587473
5008826	1	4	0.03455723542116631
5008826	1	5	0.04319654427645788
5008826	1	6	0.05183585313174946
5008826	1	7	0.06047516198704104
5008826	1	8	0.06911447084233262
5008826	1	9	0.07775377969762419
5008826	1	10	0.08639308855291576

only showing top 10 rows


```
In [41]: spark.sql("""
SELECT
    AGE_YEARS,
    SUM(BAD_FLAG) AS total_bad,
    COUNT(*) AS total_customers,
    ROUND(SUM(BAD_FLAG)/COUNT(*)*100,2) AS bad_rate
FROM credit_analysis_view
GROUP BY AGE_YEARS
ORDER BY AGE_YEARS
""").show()
```

AGE_YEARS	total_bad	total_customers	bad_rate
20	0	1	0.0
21	2	89	2.25
22	23	1597	1.44
23	41	2288	1.79
24	178	5672	3.14
25	159	7800	2.04
26	199	9992	1.99
27	458	21802	2.1
28	354	21688	1.63
29	322	18939	1.7
30	399	20944	1.91
31	266	19987	1.33
32	482	23311	2.07
33	384	22040	1.74
34	437	22310	1.96
35	351	20775	1.69
36	207	20473	1.01
37	415	25792	1.61
38	294	23176	1.27
39	271	24772	1.09

only showing top 20 rows

```
In [42]: from pyspark.sql.window import Window
from pyspark.sql.functions import lag

windowSpec = Window.partitionBy("ID").orderBy("MONTHS_BALANCE")
joined_df = joined_df.withColumn("prev_status", lag("STATUS").over(windowSpec))
joined_df.select("ID", "MONTHS_BALANCE", "STATUS", "prev_status").show(10)
```

ID	MONTHS_BALANCE	STATUS	prev_status
5009033	-16	0	null
5009033	-15	0	0
5009033	-14	0	0
5009033	-13	0	0
5009033	-12	X	0
5009033	-11	X	X
5009033	-10	X	X
5009033	-9	X	X
5009033	-8	X	X
5009033	-7	X	X

only showing top 10 rows

4. Visualization Layer (Plotly or Power BI)

Goal:

Upload Data in Mysql or NoSQL Read Dataset using Python or Pyspark and Build a visualization layer for business stakeholders to monitor key KPIs and trends using Python, Pyspark, Plotly or PowerBI (Optional).

Steps:

- **Analyse the data and perform EDA to find hidden insights.**

Technologies:

Plotly / Power BI (Mysql or MongoDB for Storage)

Outputs:

```

# 3 Load CSV files (update paths if needed)
app_pdf = pd.read_csv("/home/hadoop/application_record.csv")
cred_pdf = pd.read_csv("/home/hadoop/credit_record.csv")

# -----
# 4 EDA & Visualizations
# -----

# 4a. Income Distribution
fig1 = px.histogram(app_pdf, x="AMT_INCOME_TOTAL", nbins=50,
                    title="Income Distribution of Applicants")
fig1.show()

# 4b. Credit Status Over Months
status_count = cred_pdf.groupby('MONTHS_BALANCE')['STATUS'].count().reset_index()
fig2 = px.line(status_count, x='MONTHS_BALANCE', y='STATUS',
               title="Credit Status Count Over Months")
fig2.show()



# 4c. Correlation Heatmap of Numeric Features
numeric_cols = app_pdf.select_dtypes(include=np.number).columns
corr = app_pdf[numeric_cols].corr()
fig3 = ff.create_annotated_heatmap(
    z=corr.values,
    x=list(corr.columns),
    y=list(corr.columns),
    colorscale='Viridis'
)
fig3.show()


# 4d. Count of Applicants by Gender
gender_count = app_pdf.groupby('CODE_GENDER')['ID'].count().reset_index()
fig4 = px.bar(gender_count, x='CODE_GENDER', y='ID', title="Applicants by Gender")
fig4.show()

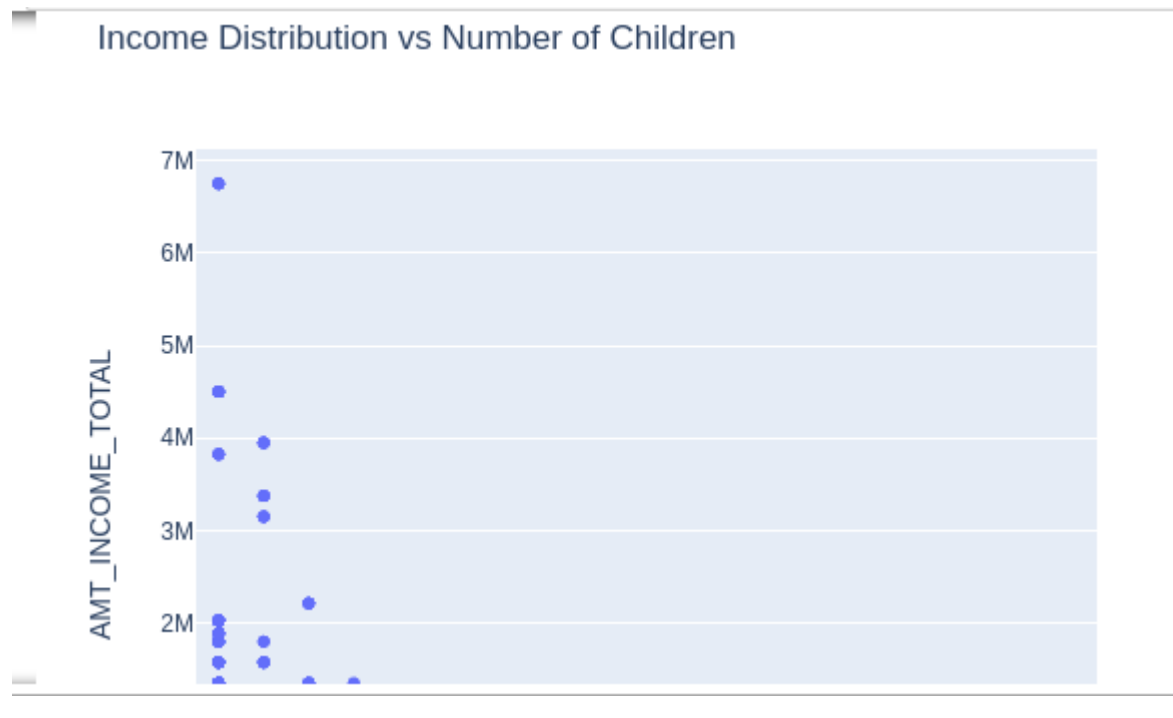
# 4e. Children vs Income
fig5 = px.box(app_pdf, x='CNT_CHILDREN', y='AMT_INCOME_TOTAL',
              title="Income Distribution vs Number of Children")
fig5.show()

```

Income Distribution of Applicants







5. Orchestration Layer (Apache Airflow)

Goal:

Automate and monitor the pipeline stages using a Directed Acyclic Graph (DAG).

Steps:

- Set up an Airflow DAG to manage the following tasks:
 - Simulate/trigger Kafka ingestion (PythonOperator / BashOperator)
 - Submit Pyspark batch job
- Schedule the pipeline to run daily/hourly based on business need.
- Add dependency handling and branching for fault tolerance.

Technologies:

Apache Airflow, PythonOperator, BashOperator

Outputs:

